

An Introduction to C++ Lambdas

August 2024

Seth Bromberger



LLNL-PRES-867878

This work was performed under the auspices of the U.S. Department of Energy by Lawrence Livermore National Laboratory under contract DE-AC52-07NA27344. Lawrence Livermore National Security, LLC



**Lawrence Livermore
National Laboratory**

Disclaimer

This document was prepared as an account of work sponsored by an agency of the United States government. Neither the United States government nor Lawrence Livermore National Security, LLC, nor any of their employees makes any warranty, expressed or implied, or assumes any legal liability or responsibility for the accuracy, completeness, or usefulness of any information, apparatus, product, or process disclosed, or represents that its use would not infringe privately owned rights. Reference herein to any specific commercial product, process, or service by trade name, trademark, manufacturer, or otherwise does not necessarily constitute or imply its endorsement, recommendation, or favoring by the United States government or Lawrence Livermore National Security, LLC. The views and opinions of authors expressed herein do not necessarily state or reflect those of the United States government or Lawrence Livermore National Security, LLC, and shall not be used for advertising or product endorsement purposes.



A gentle introduction to lambdas

- A lambda expression is an anonymous function that can inherit state from its surroundings.
- With a few (language-specific) exceptions, anything you can do in a normal function can be done with a lambda expression.
- A C++ lambda expression is an object that can be assigned to a variable or passed as an argument.
- Lambda expressions are important because they provide a concise way to control another function's behavior.
- If you've done functional programming, you've undoubtedly used lambda expressions.



Higher-order functions

Consider a function that adds two to its input:

```
int add_two(int x) { return x + 2; }  
int five = add_two(3); // 5
```

This function has a name, `add_two`. You can pass this function as an argument to another function:

```
// executor takes an int and any unary function \\  
// that has an int → int signature.  
template <typename Function>  
int executor(int n, Function fn) {  
    return fn(n);  
}  
int result = executor(3, add_two); // 5
```

`executor` is a *higher-order* function.



Lambdas: Anonymous functions

```
int add_two(int x) { return x + 2; }  
    template <typename Function>  
int executor(int n, Function fn) {  
    return fn(n);  
}  
int result = executor(3, add_two) // 5
```

If you only need `add_two` as an input to `executor`, then there is no need to assign a name to it. You can pass the function in as a *lambda*:

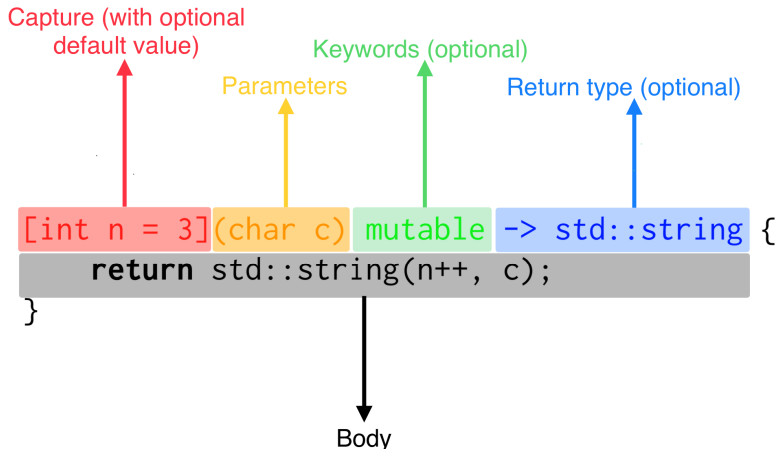
```
auto add_two { [](int x) -> int { return x + 2; } };  
int result2 = executor(3, add_two); // 5
```

or, inline:

```
int result2 = executor(3, [](int x) -> int {  
    return x + 2;  
}); // 5
```



Lambda syntax in C++



Captures

by copy

Because lambdas can be defined inline, they can *capture* the surrounding environment, bringing existing variables into the lambda body's scope:

```
int f(int x) {  
    int y = 1;  
    return executor(2, [y](int x) -> int {  
        return 2 + x + y;  
    }) // 5  
}
```



Captures

by reference

In the previous example, the lambda body has access to a *copy* of *y*. It is also possible to pass *y* by *reference*:

```
int f(int x) {  
    int y = 1;  
    return executor(2, [&y](int x) -> int {  
        return 2 + x + y++;  
    }) // 5, but now y == 2  
}
```



Mutable Captures

```
int main() {  
    int health = 3;  
    auto hit{[health]() mutable -> void {  
        health--;  
        std::cout << "Ouch! Health now " << health << "!\n";  
    }};  
    std::cout << "Health is " << health << ".\n";  
    hit();  
    hit();  
    std::cout << "Health is " << health << ".\n";  
}
```

```
Health is 3.  
Ouch! Health now 2!  
Ouch! Health now 1!  
Health is 3.
```



More Captures

There are two special capture designations that can help if you need to capture lots of variables:

- = captures everything in scope *by value*
- & captures everything in scope *by reference*
- These can be combined with specific captures.

Example:

```
int y = 5;  
int z = 10;  
    // this captures y by reference and everything else by value.  
int r = executor(2, [&y, =](int x) -> int { return x + y + z; y += 2; } ) // r=17, y=7
```

Note: it is better to be explicit about captured variables instead of relying on these “wildcard” values:

```
int r = executor(2, [&y, z](int x) -> int { return x + y + z; y += 2; } ) // r=17, y=7
```



Capture Persistence

Values of captured variables that are copied persist:

```
int main() {  
    int health = 3;  
    auto hit{[health]() mutable -> void {  
        health--;  
        std::cout << "Ouch! Health now " << health << "!\n";  
    }};  
    std::cout << "Health is " << health << ".\n";  
    hit();  
    hit();  
    std::cout << "Health is " << health << ".\n";  
    hit();  
}
```

```
Health is 3.  
Ouch! Health now 2!  
Ouch! Health now 1!  
Health is 3.  
Ouch! Health now 0!
```



Captures in YGM

In YGM, the ability to use captures depends on the function:

- In non-async functions, there are no restrictions.
- In async functions, captures are limited to *copies* of data structures that are *standard layout*¹ and *trivially copyable*²; e.g., bool, char, int, float, double, and structs of these types.
- In many scenarios (e.g., all async functions and `for_all()`), the lambda's return value is ignored and should not be specified.

¹`std::is_standard_layout`

²`std::is_trivially_copyable`



Examples

Pass by copy, no rebinding

```
int main() {  
    char c = '.';  
    int n = 3;  
    auto repeat{[c, n]() -> void { // c and n are copied.  
        std::cout << std::string(n, c) << "\n"; }  
    };  
  
    repeat();                // output: ...  
    c = '+';                 // this doesn't affect repeat  
    repeat();                // output: ...  
    [c, n]() -> void {  
        std::cout << std::string(n, c) << "\n";  
    }();                     // output: +++  
}
```



Examples

Pass by reference

```
int main() {  
    char c = '.';  
    int n = 3;  
    auto repeat{[&c, &n]() -> void {  
        std::cout << std::string(n, c) << "\n";  
    }};  
  
    repeat();           // output: ...  
    c = '+';           // wait for it!  
    repeat();           // output: +++  
    [&c, &n]() -> void {  
        std::cout << std::string(n, c) << "\n";  
    };                 // output: +++  
}
```



Examples

Another pass by reference

```
int main() {  
    char c = '.';  
    int n = 3;  
    auto repeat{[&c, &n]() -> void {  
        std::cout << std::string(n, c) << "\n"; }  
        n += 3;  
    };  
  
    repeat();                // output: ...  
    c = '+';  
    repeat();                // output: ++++++  
    [&c, &n]() -> void {  
        std::cout << std::string(n, c) << "\n";  
    };                      // output: ++++++++  
}
```



In the beginning, there was the *functor*:

```
struct add_x {  
    add_x(int n): x(n) {};           // constructor  
    int operator()(int y) { return x + y; } // override  
private:  
    int x;  
};  
  
auto add_2 = add_x(2);               // this is a struct  
int z = add_2(10);                   // 12
```

Functors capture state via the constructor, and are callable by overriding the `operator()` method.

Functors vs Lambdas

	C++ Lambda	C++ Functor
Can recurse	-	✓
Requires explicit captures	-	✓
More explicit behavior	-	✓
One-liners	✓	-
Optional return types	✓	-



